

CDC 500 Cities: Healthcare Access, Behaviors, and Health Outcomes

Stat 198 Final Project

Maya Ghanem and Isabelle Xiong

11/1/2021

Description of Data

(Include description of how you edited the data)

Research Questions

- 1) Do cities with a greater lack of healthcare access have poorer mental health and/or physical health outcomes?
- 2) Does healthcare access, mental health, and/or physical health outcomes vary by state?

Variables of Interest

Explanatory Variables:

- 1) Healthcare Access for Adults (18+): Percent of City Population that Lacks Insurance, Percent of City Population with visits to doctor for routine checkup within the past year, Percent of City Population who have high blood pressure and are taking medicine for high blood pressure control.
- 2) Geographic Distribution by State

Response Variables:

- 1) Behavior for Adults (18+): Percent of city population currently smoking, percent of city population currently reporting binge drinking habits, percent of city population reporting No leisure-time physical activity
- 2) Health Outcomes for Adults (18+): Percent of city population with coronary heart disease, percent of population diagnosed with diabetes, percent of city population with kidney disease

Linear Regressions

NOTE: Create regressions first between the explanatory (access) variables– this can indicate what kind of interactions are needed.

→ insurance vs. visits to doctor → insurance vs. medicine → visits to doctor vs. medicine

NOTE: We will not do third order interactions because they are beyond the scope of this course

Regressions for Healthcare Access and Behaviors Variables

Fit without Interaction Variables

```
access_smoking_fit <- linear_reg() %>%  
  set_engine("lm") %>%  
  fit(smoking ~ insurance + visits_to_doctor + medicine_high_bp, data = data_500_cities)  
access_smoking_fit_aug <- augment(access_smoking_fit$fit)  
tidy(access_smoking_fit) %>%  
  print()
```


```
## # A tibble: 4 x 5  
##   term                estimate std.error statistic  p.value  
##   <chr>              <dbl>    <dbl>    <dbl>   <dbl>  
## 1 (Intercept)      -15.0      2.08     -7.23 1.99e-12  
## 2 insurance         0.0523    0.0237     2.21 2.79e- 2  
## 3 visits_to_doctor -0.0966    0.0446     -2.17 3.08e- 2  
## 4 medicine_high_bp  0.674     0.0438    15.4 1.59e-43
```

```
glance(access_smoking_fit)$adj.r.squared %>%  
  print()
```


```
## [1] 0.5150724
```

```
access_smoking_fit_aug %>%  
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +  
  geom_point() +  
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +  
  labs(  
    title = "Residuals vs. Predicted City Percentage of Smoking Adults",  
    subtitle = "Data From CDC 500 Cities",  
    x = "Predicted Percentage of Smokers",  
    y = "Standard Residuals"  
  )
```

Residuals vs. Predicted City Percentage of Smoking Adults
Data From CDC 500 Cities



```
data_500_cities %>%  
  ggplot( mapping = aes(x = insurance, y = smoking)) +  
  geom_point(size = 0.25) +  
  geom_smooth(method = "lm", data = access_smoking_fit_aug, mapping = aes(x = insurance, y = .fitted))
```



```
access_binge_drinking_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(binge_drinking ~ insurance + visits_to_doctor + medicine_high_bp, data = data_500_cities)
access_binge_drinking_fit_aug <- augment(access_binge_drinking_fit$fit)
tidy(access_binge_drinking_fit) %>%
  print()
```

2) Access Variables vs. Binge Drinking

```
## # A tibble: 4 x 5
##   term          estimate std.error statistic  p.value
##   <chr>         <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)    24.2      1.58     15.3 2.65e-43
## 2 insurance     -0.162    0.0179    -9.02 4.74e-18
## 3 visits_to_doctor 0.0565   0.0337     1.68 9.45e- 2
## 4 medicine_high_bp -0.137    0.0331    -4.13 4.39e- 5
```


```
glance(access_binge_drinking_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.2367489
```


```
access_binge_drinking_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
```

```
labs(
  title = "Residuals vs. Predicted Percentage of City Reporting Binge Drinking",
  subtitle = "Data From CDC 500 Cities",
  x = "Predicted Percentage of Adults Reporting Binge Drinking",
  y = "Standard Residuals"
)
```

Residuals vs. Predicted Percentage of City Reporting Binge Drinking
Data From CDC 500 Cities



```
data_500_cities %>%
  ggplot(mapping = aes(x = insurance, y = binge_drinking)) +
  geom_point(size = 0.25) +
  geom_smooth(data = access_binge_drinking_fit_aug, mapping = aes(x = insurance, y = .fitted))
```



```
access_physical_activity_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(physical_activity ~ insurance + visits_to_doctor + medicine_high_bp, data = data_500_cities)
access_physical_activity_fit_aug <- augment(access_physical_activity_fit$fit)
tidy(access_physical_activity_fit) %>%
  print()
```

3) Access Variables vs. Physical Activity


```
## # A tibble: 4 x 5
##   term          estimate std.error statistic  p.value
##   <chr>         <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)  -28.1      1.77    -15.9  5.76e-46
## 2 insurance      0.533     0.0201    26.5  3.31e-95
## 3 visits_to_doctor 0.0625    0.0378     1.65  9.95e- 2
## 4 medicine_high_bp 0.738     0.0371    19.9  3.54e-64
```

```
glance(access_physical_activity_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.8369087
```

```
access_physical_activity_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
```

```
labs(
  title = "Residuals vs. Predicted Percentage of City Reporting No Physical Activity",
  subtitle = "Data From CDC 500 Cities",
  x = "Predicted Percentage Reporting No Lesisure Time Physical Activity",
  y = "Standard Residuals"
)
```



Fit with Interaction Variables

```
int_access_smoking_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(smoking ~ insurance + visits_to_doctor + medicine_high_bp + (insurance * visits_to_doctor) + (insurance * medicine_high_bp))
int_access_smoking_fit_aug <- augment(int_access_smoking_fit$fit)
tidy(int_access_smoking_fit) %>%
  print()
```

1) Access Variables vs. Smoking


```
## # A tibble: 7 x 5
##   term                                estimate std.error statistic  p.value
##   <chr>                                <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)                        88.9      24.0         3.70 2.41e- 4
## 2 insurance                           0.872     0.417         2.09 3.71e- 2
## 3 visits_to_doctor                    -2.13     0.362        -5.90 6.95e- 9
## 4 medicine_high_bp                    -0.756     0.463        -1.63 1.03e- 1
## 5 insurance:visits_to_doctor           0.0227    0.00634         3.59 3.69e- 4
```

```
## 6 insurance:medicine_high_bp      -0.0414    0.00628    -6.58 1.25e-10
## 7 visits_to_doctor:medicine_high_bp  0.0299    0.00667     4.48 9.60e- 6


glance(int_access_smoking_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.5691301
```

```
int_access_smoking_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
  labs(
    title = "Residuals vs. Predicted City Percentage of Smoking Adults",
    subtitle = "Data From CDC 500 Cities",
    x = "Predicted Percentage of Smokers",
    y = "Standard Residuals"
  )
```



```
data_500_cities %>%
  ggplot(mapping = aes(x = insurance, y = smoking)) +
  geom_point(size = 0.25) +
  geom_smooth(method = "lm", data = int_access_smoking_fit_aug, mapping = aes(x = insurance, y = .fitted))
```

```
int_access_binge_drinking_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(binge_drinking ~ insurance + visits_to_doctor + medicine_high_bp + (insurance * visits_to_doctor))
int_access_binge_drinking_fit_aug <- augment(int_access_binge_drinking_fit$fit)
tidy(int_access_binge_drinking_fit) %>%
  print()
```

2) Access Variables vs. Binge Drinking

```
## # A tibble: 7 x 5
##   term                estimate std.error statistic  p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)       -132.      17.8     -7.40 6.26e-13
## 2 insurance          -0.125     0.309    -0.406 6.85e- 1
## 3 visits_to_doctor    2.41      0.268     8.98 6.70e-18
## 4 medicine_high_bp    2.54      0.344     7.38 7.12e-13
## 5 insurance:visits_to_doctor -0.00655 0.00470   -1.39 1.64e- 1
## 6 insurance:medicine_high_bp  0.00686 0.00466    1.47 1.42e- 1
## 7 visits_to_doctor:medicine_high_bp -0.0401 0.00495   -8.10 4.93e-15


glance(int_access_binge_drinking_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.3488416
```

```
int_access_binge_drinking_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
  labs(
    title = "Residuals vs. Predicted Percentage of City Reporting Binge Drinking",
    subtitle = "Data From CDC 500 Cities",
    x = "Predicted Percentage of Adults Reporting Binge Drinking",
    y = "Standard Residuals"
  )
)
```

Residuals vs. Predicted Percentage of City Reporting Binge Drinking

Data From CDC 500 Cities



Why is this graph not working? Ask at some point

```
int_access_physical_activity_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(physical_activity ~ insurance + visits_to_doctor + medicine_high_bp + (insurance * visits_to_doctor))
int_access_physical_activity_fit_aug <- augment(int_access_physical_activity_fit$fit)
tidy(int_access_physical_activity_fit) %>%
  print()
```

3) Access Variables vs. Physical Activity

```
## # A tibble: 7 x 5
##   term                estimate std.error statistic    p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)        55.1      20.8      2.64  0.00845
```


```
## 2 insurance                1.96      0.361      5.42 0.0000000972
## 3 visits_to_doctor        -1.47      0.313     -4.69 0.00000361
## 4 medicine_high_bp        -0.744     0.402     -1.85 0.0646
## 5 insurance:visits_to_doctor 0.000790 0.00549     0.144 0.886
## 6 insurance:medicine_high_bp -0.0257 0.00545    -4.72 0.00000317
## 7 visits_to_doctor:medicine_high_bp 0.0271 0.00578     4.68 0.00000373
```

```
glance(int_access_physical_activity_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.8488063
```

```
int_access_physical_activity_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
  labs(
    title = "Residuals vs. Predicted Percentage of City Reporting No Physical Activity",
    subtitle = "Data From CDC 500 Cities",
    x = "Predicted Percentage Reporting No Lesisure Time Physical Activity",
    y = "Standard Residuals"
  )
```

Residuals vs. Predicted Percentage of City Reporting No Physical Activity
Data From CDC 500 Cities



Regressions for Healthcare Access and Health Outcomes

Fit with Interaction Variables

```
int_access_heart_disease_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(heart_disease ~ insurance + visits_to_doctor + medicine_high_bp + (insurance * visits_to_doctor))
int_access_heart_disease_fit_aug <- augment(int_access_heart_disease_fit$fit)
tidy(int_access_heart_disease_fit) %>%
  print()
```

4) Access Variables vs. Heart Disease

```
## # A tibble: 7 x 5
##   term                                estimate std.error statistic  p.value
##   <chr>                                <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)                        23.9        4.94        4.84 1.74e- 6
## 2 insurance                           0.352       0.0857        4.10 4.79e- 5
## 3 visits_to_doctor                    -0.480      0.0743       -6.46 2.70e-10
## 4 medicine_high_bp                    -0.289      0.0952       -3.04 2.52e- 3
## 5 insurance:visits_to_doctor           0.00239     0.00130        1.84 6.67e- 2
## 6 insurance:medicine_high_bp          -0.00780     0.00129       -6.04 3.19e- 9
## 7 visits_to_doctor:medicine_high_bp   0.00767     0.00137        5.59 3.80e- 8
```


```
glance(int_access_heart_disease_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.6667498
```

```
int_access_heart_disease_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
  labs(
    title = "Residuals vs. Predicted City Percentage of Adults with Coronary Heart Disease",
    subtitle = "Data From CDC 500 Cities",
    x = "Predicted Percentage of Adults with Coronary Heart Disease",
    y = "Standard Residuals"
  )
```

Residuals vs. Predicted City Percentage of Adults with Coronary Heart Disease

Data From CDC 500 Cities



NOTE: A linear regression is not fitting for this relationship because there is a significant pattern in the residual plot.

```
int_access_diabetes_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(diabetes ~ insurance + visits_to_doctor + medicine_high_bp + (insurance * visits_to_doctor) + (insurance * medicine_high_bp))
int_access_diabetes_fit_aug <- augment(int_access_diabetes_fit$fit)
tidy(int_access_diabetes_fit) %>%
  print()
```


5) Access Variables vs. Diabetes

```
## # A tibble: 7 x 5
##   term                                estimate std.error statistic  p.value
##   <chr>                                <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)                        69.9       11.4        6.12 1.97e- 9
## 2 insurance                           0.975      0.198        4.92 1.22e- 6
## 3 visits_to_doctor                    -1.07      0.172       -6.25 9.40e-10
## 4 medicine_high_bp                    -1.40      0.220       -6.36 4.72e-10
## 5 insurance:visits_to_doctor          -0.00935   0.00301     -3.10 2.03e- 3
## 6 insurance:medicine_high_bp          -0.00147   0.00299     -0.493 6.22e- 1
## 7 visits_to_doctor:medicine_high_bp   0.0230    0.00317      7.24 1.87e-12
```

```
glance(int_access_diabetes_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.7110294
```

```
int_access_diabetes_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
  labs(
    title = "Residuals vs. Predicted City Percentage of Adults with Diabetes",
    subtitle = "Data From CDC 500 Cities",
    x = "Predicted Percentage of Adults with Diabetes",
    y = "Standard Residuals"
  )
)
```



```
int_access_kidney_disease_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(kidney_disease ~ insurance + visits_to_doctor + medicine_high_bp + (insurance * visits_to_doctor))
int_access_kidney_disease_fit_aug <- augment(int_access_kidney_disease_fit$fit)
tidy(int_access_kidney_disease_fit) %>%
  print()
```

6) Access Variables vs. Kidney Disease


```
## # A tibble: 7 x 5
##   term                estimate std.error statistic  p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)       22.9      2.50     9.16 1.63e-18
## 2 insurance         0.198    0.0435    4.56 6.44e- 6
```

```
## 3 visits_to_doctor          -0.361      0.0377      -9.57  6.10e-20
## 4 medicine_high_bp         -0.372      0.0483      -7.70  8.53e-14
## 5 insurance:visits_to_doctor  0.000243  0.000661      0.368  7.13e- 1
## 6 insurance:medicine_high_bp -0.00297  0.000655     -4.53  7.40e- 6
## 7 visits_to_doctor:medicine_high_bp 0.00646  0.000696      9.28  6.23e-19
```

```
glance(int_access_kidney_disease_fit)$adj.r.squared %>%
  print()
```

```
## [1] 0.6193093
```

```
int_access_kidney_disease_fit_aug %>%
  ggplot(mapping = aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = "purple", lty = "dashed") +
  labs(
    title = "Residuals vs. Predicted City Percentage of Adults with Kidney Disease",
    subtitle = "Data From CDC 500 Cities",
    x = "Predicted Percentage of Adults with Diabetes",
    y = "Standard Residuals"
  )
```



NOTE: A linear regression is not fitting for this relationship because there is a significant pattern in the residual plot.

Regression With Most Correlated Variables

ANOVA Testing

Initial Visualizations

NOTE: Use initial visualizations to check if assumptions of ANOVA are met!

Overall Tests

```
summary(aov(insurance~StateDesc,data=data_500_cities)) %>%  
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## StateDesc    50   9260   185.20   8.487 <2e-16 ***  
## Residuals   424   9252    21.82  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(aov(visits_to_doctor~StateDesc,data=data_500_cities)) %>%  
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## StateDesc    50   8395   167.90   44.01 <2e-16 ***  
## Residuals   421   1606     3.81  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
## 3 observations deleted due to missingness
```

```
summary(aov(medicine_high_bp~StateDesc,data=data_500_cities)) %>%  
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## StateDesc    50   9541   190.82   44.25 <2e-16 ***  
## Residuals   422   1820     4.31  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
## 2 observations deleted due to missingness
```

```
summary(aov(smoking~StateDesc,data=data_500_cities)) %>%  
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## StateDesc    50   4752    95.03   9.747 <2e-16 ***  
## Residuals   420   4095     9.75  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
## 4 observations deleted due to missingness
```

```
summary(aov(binge_drinking~StateDesc,data=data_500_cities)) %>%  
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## StateDesc    50   1719    34.37   9.579 <2e-16 ***  
## Residuals   421   1511     3.59  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



```
## 3 observations deleted due to missingness
summary(aov(physical_activity~StateDesc,data=data_500_cities)) %>%
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)
## StateDesc    50  10657   213.13   10.76 <2e-16 ***
## Residuals   421   8343    19.82
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 3 observations deleted due to missingness
```

```
summary(aov(heart_disease~StateDesc,data=data_500_cities)) %>%
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)
## StateDesc    50   201.1    4.021   5.974 <2e-16 ***
## Residuals   421   283.4    0.673
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 3 observations deleted due to missingness
```

```
summary(aov(diabetes~StateDesc,data=data_500_cities)) %>%
  print()
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)
## StateDesc    50   1061    21.21   4.632 <2e-16 ***
## Residuals   421   1928     4.58
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 3 observations deleted due to missingness
```

```
summary(aov(kidney_disease~StateDesc,data=data_500_cities)) %>%
  print()
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## StateDesc    50   21.77   0.4354   2.102 4.48e-05 ***
## Residuals   422   87.41   0.2071
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 2 observations deleted due to missingness
```

It seems like pretty much all the overall tests indicate significant variance across the groups.

Step Down Tests

```
insurance_state_pair <- pairwise.t.test(data_500_cities$insurance, data_500_cities$StateDesc, p.adj = "none")
sig_ins_state_pairs <- broom::tidy(insurance_state_pair) %>%
  filter(p.value<0.05) %>%
  arrange(group1,group2)
nrow(sig_ins_state_pairs)
```

```
## [1] 0
```

```
print(sig_ins_state_pairs)
```

```
## # A tibble: 0 x 3
## # ... with 3 variables: group1 <chr>, group2 <chr>, p.value <dbl>
```

The overall ANOVA test says there is a significance but the Step Down tests show no significant pairs?

```
doctor_state_pair <- pairwise.t.test(data_500_cities$visits_to_doctor, data_500_cities$StateDesc, p.adjust.method = "none")
sig_doctor_state_pairs <- broom::tidy(doctor_state_pair) %>%
  filter(p.value<0.05) %>%
  arrange(group1,group2)
nrow(sig_doctor_state_pairs)
```

```
## [1] 0
```

```
print(sig_doctor_state_pairs)
```

```
## # A tibble: 0 x 3
```

```
## # ... with 3 variables: group1 <chr>, group2 <chr>, p.value <dbl>
```

Ok so there is clearly an issue. I think I am interpreting the F-statistic incorrectly?

```
smoking_state_pair <- pairwise.t.test(data_500_cities$smoking, data_500_cities$StateDesc, p.adjust.method = "none")
sig_smoking_state_pairs <- broom::tidy(smoking_state_pair) %>%
  filter(p.value<0.05) %>%
  arrange(group1,group2)
nrow(sig_smoking_state_pairs)
```

```
## [1] 0
```

```
print(sig_smoking_state_pairs)
```

```
## # A tibble: 0 x 3
```

```
## # ... with 3 variables: group1 <chr>, group2 <chr>, p.value <dbl>
```